

Rohit Annasaheb Ragde

Philadelphia, PA · ragderohit2025@gmail.com · +1 (267) 278-1779 · linkedin.com/in/rohitragde · github.com/rohitragde

PROFESSIONAL SUMMARY

Senior Data Engineer with 6+ years of experience architecting and delivering enterprise-scale data pipelines, cloud data platforms, and analytics infrastructure across financial services, utilities, and telecommunications. Deep expertise in PySpark, dbt, Airflow, Snowflake, and Azure Databricks — proven track record transforming high-volume raw data into reliable, analytics-ready systems that directly drive business decisions. Adept at leading cross-functional collaboration on cloud migration initiatives and delivering scalable, secure, high-performance data solutions.

TECHNICAL SKILLS

Languages: Python, SQL, T-SQL, PL/SQL, Bash, Unix Shell

Data Engineering: ETL/ELT Pipelines, Data Modeling, CDC, SCD, Data Quality & Validation, Schema Design

Big Data & Streaming: PySpark, Apache Spark, Apache Kafka (familiar), Spark Streaming

Cloud Platforms: Azure Databricks, Azure Data Lake Gen2, Azure Data Factory, AWS S3, AWS Glue (familiar), Snowflake, Teradata Vantage

Orchestration & Transform: Apache Airflow, dbt, Azure Data Factory

GenAI & AI Engineering: Claude API, Prompt Engineering, Multi-turn Conversational Workflows, API-driven AI Agent Orchestration, JSON/API Response Transformation

Databases: SQL Server, PostgreSQL, Oracle, MongoDB, Teradata

DevOps & Tools: Git, GitHub, Docker, CI/CD, Jenkins, Power BI, Tableau

Security & Compliance: TDE, RBAC, Auditing, Data Masking, GDPR, HIPAA, SOX

PROFESSIONAL EXPERIENCE

TD Bank — Financial Services Philadelphia, PA

Senior Data Engineer (Financial Data) Oct 2025 – Present

Architected and maintained scalable SQL-based data pipelines integrating 15+ source systems powering regulatory reporting and AML analytics workflows for one of Canada's largest banks

Engineered 30+ automated data quality validation frameworks using Python and SQL — null checks, reconciliation logic, anomaly detection — reducing data incidents by 40% and ensuring SOX compliance

Designed Power BI dashboards tracking 20+ financial KPIs consumed daily by finance leadership, enabling data-driven decisions across retail banking operations

Automated manual reconciliation and reporting workflows saving 40+ hours/month, improving operational efficiency and freeing analyst capacity for higher-value work

Performed large-scale variance analysis and data reconciliation across financial datasets, identifying systemic discrepancies and improving downstream reporting reliability

Cantonica (StableCoupons Inc) — Technology Startup Philadelphia, PA

Data Engineer — Data Platform & Analytics Sep 2024 – Sep 2025

Designed and built end-to-end data pipelines using Snowflake and dbt, implementing structured dimensional data models that reduced downstream data preparation time by 30%

Optimized complex SQL queries and schema design improving data processing performance by 25%, enabling real-time analytics on high-velocity product usage data

Implemented robust data quality and validation frameworks (null checks, reconciliation, monitoring alerts) ensuring pipeline reliability across production workloads

Integrated data pipelines with Django application backend using MongoDB and contributed to UI data rendering for data-driven product features

Exelon Inc — Utility Services Wilmington, DE

Data Engineer — Advanced Data Analytics Jun 2023 – Dec 2023

Built and optimized high-volume data pipelines using PySpark on Azure Databricks processing 15+ TB of daily utility operations data with 95% SLA reliability

Developed scalable PySpark transformation pipelines for 10+ datasets applying business rules, schema normalization, and data standardization across on-premises SQL Server and Azure Synapse environments

Reduced critical data load execution time by 30% by identifying bottlenecks, optimizing transformation logic, and implementing indexing strategies across 20+ enterprise tables

Resolved production data quality issues by implementing comprehensive validation frameworks using Python and SQL, reducing data anomalies by 40% and establishing automated alerting systems

Engineered Azure Data Factory workflows and Python-based transformations to ETL enterprise data from on-premises SQL Server to Azure Synapse, reducing data latency by 45%

Teradata Inc — Enterprise Data Consulting Mumbai, India

Data Integration Engineer (Clients: Audi, Bank of Ireland, Volkswagen) Sep 2019 – Aug 2022

Led client engagement initiatives developing SLIPSTREAM — an automated ELT pipeline using Teradata and Talend for telecom customer analytics, processing 10M+ subscriber records daily with cron-based scheduling, improving revenue forecasting accuracy by 25%

Architected dimensional data models and maintained 10+ enterprise ETL pipelines using Teradata, Talend, and SQL Server across retail, manufacturing, and telecom clients, processing 2.5TB daily with 95% reliability

Optimized 100+ complex SQL queries using execution plan analysis, indexing strategies, and stored procedure development, transforming query response times from minutes to seconds on enterprise SQL Server

Built 40+ automation and monitoring scripts using Python and Bash for pipeline failure detection, implementing automated recovery mechanisms that reduced mean time to resolution from hours to minutes

Developed inventory optimization and customer recommendation systems processing 20M+ retail transactions using Teradata and Python, delivering dynamic pricing solutions that increased client revenue by 15%

PROJECTS

AI Agent Orchestration System — Personal Project

GenAI + Data Engineering Mar 2026 – Present

Engineered a multi-agent AI system using Claude API featuring 6 specialized agents with multi-turn conversational workflows, real-time data retrieval, and automated pipeline orchestration

Implemented prompt engineering patterns, JSON/API response transformation, and metadata-driven architecture integrating 4 external APIs — Anthropic, Happenstance, Gmail, Greenhouse — with automated fallback logic

Built API-driven AI workflows with structured error handling, SQLite persistence layer, and Windows Task Scheduler automation — demonstrating production-grade GenAI system design

Scalable Financial Transactions Data Pipeline

PySpark + dbt + Airflow Mar 2026 – Apr 2026

Designed end-to-end pipeline ingesting and processing financial transaction data simulating real-world banking analytics — fact/dimension modeling (transactions, users, merchants) with PySpark transformations and dbt

Orchestrated pipeline execution using Airflow with automated validation frameworks (null handling, deduplication, anomaly detection) ensuring reliable financial reporting

NYC Taxi Analytics Pipeline — Medallion Architecture

Apache Spark + Delta Lake Feb 2026 – Mar 2026

Built Bronze/Silver/Gold medallion architecture processing high-volume NYC taxi datasets with Spark-based transformations, partitioning strategies, and incremental data processing

Implemented schema enforcement, data quality monitoring, and anomaly detection enabling business insights into trip patterns, peak demand, and revenue drivers

EDUCATION

Drexel University — Philadelphia, PA Aug 2022 – Jun 2024

Master of Science, Computer Science

University of Mumbai — Mumbai, India Jul 2016 – May 2019

CERTIFICATIONS

Teradata Database Vantage Certified Administrator — Teradata Jun 2021

Teradata Database Vantage Certified SQL Associate — Teradata Mar 2021

Oracle Database SQL 1Z0-071 Certified SQL Associate — Oracle Aug 2018